

Tello Control — Technical Documentation

Gesture- and voice-controlled DJI Tello, fully local, hardware-optional

Marcelo Leindl

2026

Contents

1. Overview & motivation	1
2. System architecture	2
2.1 The backend abstraction	2
2.2 Data flow	2
3. Component deep-dives	2
3.1 Core — <code>tello_control.core</code>	2
3.2 Gesture control — <code>tello_control.gesture</code>	3
3.3 Voice control — <code>tello_control.voice</code>	3
3.4 Simulation — <code>tello_control.sim</code>	4
4. Technology stack	4
5. Setup & run	5
6. Engineering decisions & lessons	5
6.1 Explored and ruled out — lip-reading control	5
7. Status & roadmap	7

1. Overview & motivation

Tello Control lets a DJI Tello drone be flown by **hand gestures** and **spoken commands**, with everything running **locally on the laptop — no cloud, no internet during flight**. The defining design idea is a *hardware-optional* architecture: the entire control stack can be developed, tested and demonstrated **without owning or risking a real drone**, because a single abstraction can drive three interchangeable backends:

- a **software mock** (pure logic, instant, no Wi-Fi),
- a **PyBullet physics simulation** (real quadrotor dynamics in a 3D window),
- the **real Tello** over Wi-Fi.


Why this matters: drones are unforgiving to debug. A wrong sign in a movement command, a missing safety check, or a flaky gesture classifier becomes a crashed drone. By making the mock and the simulation first-class backends behind the *same* interface, the risky parts (gesture logic, language parsing, command validation) are hardened off-hardware first, and only then pointed at the real drone by changing one argument.

This document explains **what each component does, which libraries are used, and how they connect**.

2. System architecture

2.1 The backend abstraction

Everything above the controller is backend-agnostic:



`DroneController` (`tello_control/core/controller.py`) exposes a small, stable verb set — `connect`, `takeoff`, `land`, `emergency`, `forward/back/left/right/up/down`, `rotate_cw/ccw`, `disconnect` — and forwards each call to the chosen backend. The gesture and voice modules only ever see this interface; they never import a concrete drone class. Switching from simulation to a real flight is literally one argument: `DroneController(backend="sim")` → `DroneController(backend="real")`.

The real-drone class (`djitellopy.Tello`) is imported **lazily**, only when `backend="real"` is requested, so the package runs fine on a machine that never connects to a drone. The same is true for the `PyBullet` backend, which lives in a separate conda environment.

2.2 Data flow

Both input modalities converge on the same controller:

- **Gesture:** webcam frame → MediaPipe hand landmarks → angle-based classifier → a `Gesture` enum → `debounce` → a `DroneController` method call.
- **Voice:** microphone → voice-activity detection + wake word → speech-to-text → local LLM → JSON command list → validation → a sequence of `DroneController` calls.

3. Component deep-dives

3.1 Core — `tello_control.core`

`MockTello` is a software stand-in for `djitellopy.Tello`. It mirrors the real SDK’s method names but, instead of flying, it maintains a simulated 3-D pose (`x`, `y`, `z` in cm and `yaw` in degrees), logs every command with a timestamp, and enforces the real SDK’s limits (distance 20–500 cm, angle 1–360°) by raising a `TelloException` on violations. It can print a command log and an ASCII top-down map of the flight path. This is what makes “test the whole control logic at your desk” possible.

`DroneController` is the abstraction layer described above. It also offers `position()` and `report()` helpers that are only meaningful when a pose is tracked (mock and sim), and a `tick()` method that advances the simulation physics one step (a no-op for mock/real).

The inheritance trick: the simulation backend (`PyBulletBackend`) *subclasses* `MockTello`. It therefore inherits the entire logic layer — pose tracking, logging, the map, and the SDK safety checks — for free. It overrides only the flight methods, which first call `super()` (analytic pose update + logging + checks, exactly

like the mock) and *then* drive the real physics toward that pose. The command semantics stay identical to the mock; the simulation just adds real flight dynamics on top.

3.2 Gesture control — `tello_control.gesture`

Pipeline: **webcam** → **MediaPipe** → **classifier** → **debounce** → **command (async)**.

- `detector.py` runs MediaPipe's `HandLandmarker` (Tasks API) on each webcam frame and classifies the hand. Classification is **angle-based, not position-based**: a finger counts as extended when the joint angle at its PIP joint (MCP → PIP → TIP) is large (~straight). This is robust to hand rotation, tilt and size, which a position-based check is not. Thumb direction is read from the thumb tip relative to the knuckle line; pointing direction uses the index finger's angle to the vertical with a $\pm 30^\circ$ dead zone for "forward".
- `command_map.py` maps each `Gesture` to a controller call, with a **debounce**: a gesture must be held for `STABLE_FRAMES` (5) consecutive frames before it fires, and a `COOLDOWN_FRAMES` (10) pause then blocks repeats. This prevents jitter and accidental double-triggers. Both were halved from their original values (8 / 20) to cut the gesture-to-command latency.
- `runner.py` runs flight commands on a **worker thread** (`AsyncCommandRunner`) so the blocking flight call doesn't freeze the webcam/preview loop. A `ThreadedCtrlAdapter` makes the async runner look like a normal controller, so the mapping code stays unchanged. While a command runs, new ones are dropped (flight commands must not stack up).
- `app.py` is the main loop: it opens the webcam, draws an overlay (current gesture, a debounce progress bar, drone state), advances the sim physics each frame via `tick()`, and handles keys (`t` takeoff, `e` emergency, `q` quit).

Gesture map: fist → hover, thumb up/down → up/down, index up/left/right → forward/left/right, peace → back, open hand → land.

Continuous velocity mode (--rc, opt-in). The default path sends discrete 30 cm `move_*` commands, each of which blocks until the drone acknowledges it — roughly 200 ms of dead time per command. `velocity_map.py` offers the alternative: a *held* gesture is translated into an RC setpoint and pushed with `send_rc_control`, which is fire-and-forget. Motion begins the moment the gesture is recognised.

`VelocityBlender` ramps the output toward the target by at most `MAX_STEP` per axis per frame. That smooths acceleration and braking, and it makes the frame debounce unnecessary as a side effect: a single misclassified frame nudges the setpoint by one step and is pulled straight back, so it never becomes visible motion. A dead zone suppresses jitter near zero. `--rc --sim` deliberately raises — driving a velocity setpoint through the PyBullet PID controller is a separate piece of work.

3.3 Voice control — `tello_control.voice`

Pipeline: **mic** → **VAD + wake word** → **Whisper** → **Ollama** → **JSON** → **validation** → **execute**.

- `stt.py` records 16 kHz mono audio (`sounddevice`) and transcribes German speech with **faster-whisper** (small model, int8, CPU). The model is loaded once and reused.
- `listener.py` powers the always-listening mode. A pure-numpy **energy VAD** (RMS threshold, auto-calibrated to the room's noise floor) cuts the mic stream into speech segments. A command is only forwarded if the transcript starts with the **wake word "Drohne"**, which is then stripped — this is the safety gate in continuous mode.
- `llm_parser.py` sends the transcript to a **local LLM via Ollama** (qwen2.5:3b, JSON mode, temperature 0) with a German system prompt and few-shot examples (including tricky number words like "ein halber Meter" = 50 cm). It also has an autostart helper (`ensure_ollama`) that launches the Ollama server if needed and checks the model is pulled. The Ollama call is deliberately separated from parsing so the parse/validate logic is testable without a running server.
- `fastpath.py` short-circuits the LLM for unambiguous single-clause commands ("lande", "vor 50 cm", "dreh dich rechts"). Anchored regexes map the transcript straight onto the command schema, saving

the ~0.8 s Ollama round-trip. Anything with conjunctions, or any value outside the SDK bounds, returns `None` and falls through to `llm_parser.parse()`. Crucially the fastpath emits its result through the **same** `validate_list()` as the LLM path — it is a latency shortcut, never a safety bypass.

- `commands.py` is the **safety-critical validation layer**. The LLM returns a JSON command list; every command is checked against an allow-list of actions and the SDK bounds *before anything reaches the drone*. The contract is all-or-nothing: if any command is invalid, **nothing** executes — no half-dangerous partial flight.
- `app.py` ties it together with two modes: an ENTER-to-speak mode (with a confirmation step before real flights) and the continuous wake-word mode.

3.4 Simulation — `tello_control.sim`

Runs in the separate conda environment `tello-sim` (see §5).

- `pybullet_backend.py` wraps **gym-pybullet-drones** (built on PyBullet). On `connect()` it spawns a `CtrlAViary` with a Crazyflie model and a `DSLIPIDControl` PID controller, opens a 3-D window, and (optionally) follows the drone with the camera. Coordinates are converted between the MockTello convention (cm, clockwise yaw) and the PyBullet world (m, counter-clockwise yaw). It supports a *cooperative* mode where flight commands don't block: they only set the target pose, and the physics is advanced one slice per frame via `tick()` from the main thread — needed for interactive gesture/voice loops, because PyBullet's GUI must run on the main thread.
- `control_lab.py` is a small **control-engineering lab**: it flies a 1-metre step in X with different proportional gains and measures overshoot and settling time, saving a plot to `results/control_lab_step.png`. This is something the real Tello can *not* do — the Tello SDK only exposes velocity setpoints, not motor-level access — which is exactly why it lives in the simulation.
- `demo.py` flies the cube demo in the 3-D window; `launcher.py` is a menu that starts each sim program in its own process (so OpenCV and PyAV, which both bundle ffmpeg, never collide in one process).

4. Technology stack

Layer	Library / tool	Role	Why
Real drone	djitellopy	Tello SDK over Wi-Fi (UDP)	de-facto Python Tello SDK
Logic mock	MockTello (own)	hardware-free drone	test logic at the desk
Physics sim	PyBullet / gym-pybullet-drones	real quadrotor dynamics + PID	native on Apple Silicon, no Docker; ready-made controllers
Gestures	MediaPipe	hand landmarks	fast, local, robust hand tracking
Vision I/O	OpenCV	webcam capture + overlay	standard
Speech ↔ text	faster-whisper	local STT	accurate German STT, fully offline
Audio I/O	sounddevice	mic capture	simple numpy-native streaming

Layer	Library / tool	Role	Why
Text $\rightarrow$ command	Ollama (qwen2.5:3b)	local LLM JSON parsing	offline, small, JSON mode
Controller input	pygame	PS4 controller	manual flight option
Tests	pytest	hardware-free test suite	guards logic + safety

5. Setup & run

Two environments are used:

1. **Main** (venv, Python 3.11+): mock, real drone, gestures, voice. `pip install -e .` then `python scripts/download_model.py`.
2. **Simulation** (conda env `tello-sim`): PyBullet. Created from `environment-sim.yml`. PyBullet is installed via **conda-forge**, **not pip**, because on Apple Silicon the pip build of PyBullet fails (its bundled zlib uses old C that recent Apple clang rejects).

Full command list: see [COMMANDS.md](#).

6. Engineering decisions & lessons

- **Gazebo $\rightarrow$ PyBullet.** Gazebo on Apple Silicon meant Docker and a painful GUI; PyBullet runs natively, gives real quadrotor physics and ready-made PID controllers — a better fit for a laptop-only project.
- **conda vs pip on M1.** PyBullet’s pip build breaks on Apple clang; conda-forge ships a working binary. `gym-pybullet-drones` is installed editable with `pip install --no-deps` so it can’t pull a second PyBullet or drag in SB3 + PyTorch.
- `setuptools < 81`. `gym-pybullet-drones` still uses `pkg_resources`, removed in newer `setuptools` — pinned in the sim environment.
- **OpenCV $\rightarrow$ PyAV conflict.** Both bundle `ffmpeg`, causing duplicate-class warnings. The sim launcher therefore runs each mode in its own process.
- **Threading model for gestures.** Flight calls `block`; the webcam loop must not. A single worker thread owns the drone, the main thread owns the camera and the PyBullet GUI — no race conditions, smooth preview.
- **Validation as a safety gate.** The voice path treats the LLM as untrusted: its output is validated against an allow-list and SDK bounds, all-or-nothing, before any command reaches the drone.

6.1 Explored and ruled out — lip-reading control

A lip-driven control channel was prototyped and ultimately dropped. Two approaches were tested:

1. **Lip-shape templates (MediaPipe FaceLandmarker + nearest-neighbour).** A 30-second per-user calibration recorded one normalised 80-dim lip vector per command, matched live with debounce/cooldown — the same architecture as the gesture pipeline. It runs in real time and is fully local, but the mouth-shape vectors for distinct commands were not separable enough on a single face: short words spend most frames near the neutral (closed-mouth) shape, so matching was unreliable even with a neutral-frame filter and multi-take calibration.
2. **True lip-reading (Meta AV-HuBERT, visual speech recognition).** The full VSR pipeline was reproduced locally: webcam $\rightarrow$ `dlib` 68-point landmarks $\rightarrow$ mean-face affine alignment $\rightarrow$ 96×96 mouth ROI $\rightarrow$ centre-crop 88 $\rightarrow$ AV-HuBERT base (LRS3-433h) $\rightarrow$ beam search. Preprocessing was confirmed correct (clean landmark detection and aligned ROI sequences), and the model ran end-to-end. Accuracy, however, was unusable for control: “take off” decoded to “*land earth*”, and “please turn




Figure 1: PID step response measured in the simulation: a gentle gain ($P=0.20$), the default ($P=0.40$), and an aggressive gain ($P=0.90$) that overshoots heavily.

off and fly up” to “*but lies a lot of guidance overly helping*” — fluent, hallucinated LRS3-style English rather than the spoken words. **Root cause:** AV-HuBERT is trained on natural connected British speech; short isolated commands are its worst case, compounded by a non-native speaker off the training distribution. The ~5 s clip-buffer + inference latency also rules it out for reactive flight.

Conclusion. Speaker-independent, low-latency lip reading is not achievable off-the-shelf today; making AV-HuBERT usable would require fine-tuning on a per-user command dataset — a separate ML project. The lip-shape prototype was removed from the codebase. Gesture and voice remain the two working control channels.

7. Status & roadmap

- **A0 Foundation** □ — controller abstraction, MockTello, first real flights.
- **A1 Gesture control** □ — MediaPipe, angle-based classifier, debounce; flown on the real drone in both discrete and `--rc` velocity mode.
- **A2 Voice control** □ — Whisper □ Ollama □ validated JSON, wake word; flown on the real drone.
- **A3 Physics sim** □ — PyBullet backend + PID control lab.
- **A4 Integration** □ — unified app, portfolio polish.